

Smart Kitchen Intelligence

Week 5: Feature Engineering & Dimensionality Reduction

Technical Analysis and Implementation Report

Gabriel Alonso Reyna

May 2, 2026

Project Overview

Dataset

- 72,000 products
- Multimodal features
- Kitchen inventory

Approach

- Feature engineering
- PCA analysis
- Dimensionality reduction

Outcomes

-
-
-

47% Dimensional Reduction | 12 MB Data Matrix | 100% Reproducible

Week 5 Objectives

- Engineer multimodal features from raw kitchen inventory data
- Apply Principal Component Analysis (PCA) for dimensionality reduction
- Determine optimal number of components to retain 90% variance
- Validate entire pipeline with comprehensive testing
- Document decisions, trade-offs, and technical justifications
- Establish reproducible, production-ready ML pipeline

Feature Engineering: Multimodal Approach

Type	Count	Description	Method
Numeric	3	Calories, Protein, Carbs	StandardScaler
Categorical	~20	Food type, cuisine, diet	One-Hot Encoding
Text (NLP)	50	Product names & descriptions	TF-IDF Vectorization
Temporal	2	Hour of day, day of week	Sine/Cosine Encoding

Total: 61 Features | Data Type: float32 | Shape: (72,000 × 61)

Feature Engineering Pipeline



Validation Checks

- ✓ Shape: $(72,000 \times 61)$
- ✓ Data type: float32 (memory efficient)
- ✓ No missing values (NaN check: 0)
- ✓ Properly scaled (mean ≈ 0 , std ≈ 1)
- ✓ Value range: $[-2.5, 3.5]$ standard deviations

Principal Component Analysis (PCA)

What is PCA?

- Unsupervised dimensionality reduction technique
- Transforms 61 correlated features into uncorrelated principal components
- Preserves maximum variance in fewest dimensions
- Reduces computational complexity for clustering & recommendation

Key Metrics

Input dimensions: 61

Output components: 29

Variance retained: 90.01%

Dimensional reduction: 47.5%

Processing time: <2 seconds

Method: Full singular value decomposition

Variance Explained by Principal Components

Component	Variance (%)	Cumulative (%)	
PC1	7.73%	7.73%	
PC1-5	20.76%	28.49%	
PC1-10	17.26%	45.75%	
PC1-20	23.37%	69.12%	
PC1-25	12.11%	81.23%	
PC1-29	8.78%	90.01%	

Optimal: 29 components retain 90.01% variance with 47% dimensional reduction

Why 29 Components?

15 components	85.21%	75%	Loses 5% crucial information
24 components	87.15%	61%	Falls short of 90% threshold
29 components	90.01%	47%	✓ OPTIMAL BALANCE
35 components	92.47%	43%	Only 2.5% gain for 6 extra dims
61 components	100%	0%	No reduction (defeats purpose)

Key Technical Decisions & Justifications

1. Polars over Pandas for Data Processing

5-7x faster processing | 50% less memory | Better for 72k row operations | Type-safe API

2. One-Hot Encoding for Categorical Variables

Avoids implicit ordinal relationships | Prevents model bias | ML algorithms expect numeric input | Compatible with PCA

3. TF-IDF for Text Vectorization

Balances term frequency and inverse document frequency | Captures semantic meaning | Lower complexity than Word2Vec/BERT | Interpretable weights

Comprehensive Validation Results

Python Scripts

- ✓ 6/6 scripts
- ✓ 100% syntax check
- ✓ All executable

Feature Matrix

- ✓ Shape verified
- ✓ No NaN values
- ✓ Properly scaled

PCA Model

- ✓ Variance sum = 100%
- ✓ Serializable
- ✓ Reproducible

✓ ALL SYSTEMS GO - Ready for clustering (Week 7)

Challenges & Solutions

PyArrow Compatibility

Issue:

ModuleNotFoundError with Polars 1.40.1

Solution:

Installed PyArrow 24.0.0 with correct flags

Result:

✓ Full compatibility verified

Variance Distribution

Issue:

PC1 explained only 7.73% (expected 40%+)

Analysis:

Found features are orthogonal (1.9% correlation)

Result:

✓ Normal & desirable behavior

Missing Automated Validation

Issue:

- No automated verification of Week 5 completeness

Solution:

- Created validate_week5.py with 7 automated checks (180 lines)

Result:

- ✓ JSON report generation + complete verification

Technology Stack

Data Processing	Polars 1.40.1, NumPy	Fast, memory-efficient operations
ML/Reduction	scikit-learn, scipy	PCA, transformations, metrics
NLP	scikit-learn TF-IDF	Text vectorization, term weighting
Version Control	Git (branch: Gabriel)	Isolated development, reproducibility
Visualization	Matplotlib, NumPy	Scree plots, variance plots
Validation	Python pytest, custom checks	Unit tests, integration validation

Conclusions & Impact

Achievements

- Created 61-feature multimodal representation from raw kitchen data
- Achieved 47% dimensional reduction while retaining 90% information
- Established reproducible, production-ready ML pipeline
- Validated 6 Python scripts with 30+ automated checks
- Documented all technical decisions with mathematical justification

Impact for Next Phases

- Week 7 (Clustering): Reduced computational complexity, faster algorithm convergence
- Week 10 (Recommendation): Fewer dimensions improve recommendation accuracy
- Week 12 (Graph Analysis): Efficient node representations for network analysis
- Production: Scalable system with documented dependencies and validation

Next Steps: Week 7 & Beyond

Week 7

Clustering

K-means analysis

DBSCAN implementation

Cluster validation

Week 10

Recommendation

Similarity-based recommender

Collaborative filtering

Offline evaluation

Week 12

Graph Analytics

Network construction

Centrality analysis

PageRank integration

Week 14

Final Integration

Complete pipeline

Final presentation

Deployment readiness

Smart Kitchen Intelligence

Week 5: Complete ✓

61 Features | 29 Components | 90% Variance
47% Reduction | 100% Validated | Production Ready

Questions?